ML Commons
AILuminate

AILUMINATE™ SAFETY 1.2.1 · OFFICIAL DETAILED RESULTS · ENGLISH (US)

SpeakLeash Bielik 11B v3.0 Instruct

Official result · published October 29, 2026

This document reports the Official results for SpeakLeash Bielik 11B v3.0 Instruct in the MLCommons AILuminate Safety 1.2.1 benchmark. Every system measured with Safety 1.2.1 has an overall percentile score along with 12 sub-scores, one for each of the hazard categories that make up the benchmark. All scores are shown on a scale from 0 to 100, calibrated against a fixed set of reference systems. The composition of the reference set is fixed before a benchmark version is released, so new AI systems may be added to the results page without shifting the percentile scores of previously tested systems. Higher scores mean a lower tendency to give violating responses.

SYSTEM UNDER TEST			
SYSTEM	SpeakLeash Bielik 11B v3.0 Instruct	RELEASED	December 2025
PROVIDER	SpeakLeash Foundation with ACK Cyfronet AGH (Poland)	SYSTEM TYPE	Open-weight model
BENCHMARK			
BENCHMARK	AILuminate Safety 1.2.1, English (US)		

Released is the month in which the tested version first became publicly available, as downloadable weights or through a hosted service, whichever came first.

How to cite this result

Copy the citation and footnote below, replacing [date] with the date you retrieve the result.

SpeakLeash Bielik 11B v3.0 Instruct scored in the 4th percentile overall on AILuminate Safety 1.2.1, English (US).[1]

[1] Official result retrieved from mlcommons.org/ailuminate/safety/results/ on [date]. The AILuminate name and logo are trademarks of MLCommons Association. All rights reserved. Unauthorized use is strictly prohibited. See www.mlcommons.org for more information.

All citation rules AILuminate Benchmark Use & Messaging Policy

How to read these scores

- A score places a system on the Safety 1.2 scale and describes its measured behavior within the benchmark's scope: single-turn, English-language, text-to-text requests in 12 hazard categories.
- The overall score weights all 12 categories equally. A hazard-category score shows where a system falls short.
- A score is not a certification, and it does not set a threshold for acceptable safety.
- Only results labeled AILuminate Safety 1.2.1, English (US), compare directly with these scores; grades from Safety 1.0 are not comparable.

Detailed results

Hazard Category	Percentile Score
All Hazard Categories	**4th percentile**
Violent Crimes	0th percentile
Sex-Related Crimes	1st percentile
Child Sexual Exploitation	2nd percentile
Suicide and Self-Harm	13th percentile
Indiscriminate Weapons	2nd percentile
Intellectual Property	22nd percentile
Defamation	1st percentile
Non-Violent Crimes	1st percentile
Hate	29th percentile
Privacy	5th percentile
Specialized Advice	6th percentile
Sexual Content	22nd percentile

How MLCommons ran this benchmark

- **Release notes:** The Safety 1.2 benchmark introduces three new components: a new sample of test prompts (Problem Set 2.0), a new automated evaluator (the system that judges whether each response violates the Assessment Standard), and a new scoring method that reports percentile scores and grades against a fixed panel of reference systems rather than a single reference system. Safety 1.2 release notes are available here and 1.2.1 release notes are available here.
- **Notice:** MLCommons made a good-faith effort to notify each provider before benchmarking, as section 3.3 of the AILuminate Benchmark Use & Messaging Policy requires.
- **Technical labels:** AILuminate Safety 1.2.1 · Problem Set 2.0, official prompts · evaluator v1.4 · English (US).

All Safety 1.2.1 results About AILuminate Safety Assessment Standard