ML Commons
AILuminate

AILUMINATE™ SAFETY 1.2.1 · OFFICIAL DETAILED RESULTS · ENGLISH (US)

Moonshot AI Kimi K2.7 Code

Official result · published October 29, 2026

This document reports the Official results for Moonshot AI Kimi K2.7 Code in the MLCommons AILuminate Safety 1.2.1 benchmark. Every system measured with Safety 1.2.1 has an overall percentile score along with 12 sub-scores, one for each of the hazard categories that make up the benchmark. All scores are shown on a scale from 0 to 100, calibrated against a fixed set of reference systems. The composition of the reference set is fixed before a benchmark version is released, so new AI systems may be added to the results page without shifting the percentile scores of previously tested systems. Higher scores mean a lower tendency to give violating responses.

SYSTEM UNDER TEST			
SYSTEM	Moonshot AI Kimi K2.7 Code	RELEASED	June 2026
PROVIDER	Moonshot AI	SYSTEM TYPE	Open-weight model
BENCHMARK			
BENCHMARK	AILuminate Safety 1.2.1, English (US)		

Released is the month in which the tested version first became publicly available, as downloadable weights or through a hosted service, whichever came first.

How to cite this result

Copy the citation and footnote below, replacing [date] with the date you retrieve the result.

Moonshot AI Kimi K2.7 Code scored in the 33rd percentile overall on AILuminate Safety 1.2.1, English (US).[1]

[1] Official result retrieved from mlcommons.org/ailuminate/safety/results/ on [date]. The AILuminate name and logo are trademarks of MLCommons Association. All rights reserved. Unauthorized use is strictly prohibited. See www.mlcommons.org for more information.

All citation rules AILuminate Benchmark Use & Messaging Policy

How to read these scores

- A score places a system on the Safety 1.2 scale and describes its measured behavior within the benchmark's scope: single-turn, English-language, text-to-text requests in 12 hazard categories.
- The overall score weights all 12 categories equally. A hazard-category score shows where a system falls short.
- A score is not a certification, and it does not set a threshold for acceptable safety.
- Only results labeled AILuminate Safety 1.2.1, English (US), compare directly with these scores; grades from Safety 1.0 are not comparable.

Detailed results

Hazard Category	Percentile Score
All Hazard Categories	**33rd percentile**
Violent Crimes	24th percentile
Sex-Related Crimes	34th percentile
Child Sexual Exploitation	35th percentile
Suicide and Self-Harm	53rd percentile
Indiscriminate Weapons	20th percentile
Intellectual Property	45th percentile
Defamation	32nd percentile
Non-Violent Crimes	22nd percentile
Hate	33rd percentile
Privacy	30th percentile
Specialized Advice	52nd percentile
Sexual Content	25th percentile

How MLCommons ran this benchmark

- **Release notes:** The Safety 1.2 benchmark introduces three new components: a new sample of test prompts (Problem Set 2.0), a new automated evaluator (the system that judges whether each response violates the Assessment Standard), and a new scoring method that reports percentile scores and grades against a fixed panel of reference systems rather than a single reference system. Safety 1.2 release notes are available here and 1.2.1 release notes are available here.
- **Notice:** MLCommons made a good-faith effort to notify each provider before benchmarking, as section 3.3 of the AILuminate Benchmark Use & Messaging Policy requires.
- **Technical labels:** AILuminate Safety 1.2.1 · Problem Set 2.0, official prompts · evaluator v1.4 · English (US).

All Safety 1.2.1 results About AILuminate Safety Assessment Standard